EECS E6893 Big Data Analytics
Intro to Big Data Analytics on GCP

Frank Ou Yang

Agenda

- GCP
- Cloud Storage
- BigQuery
- Dataproc

Google Cloud Platform (GCP)

GCP

- Cloud computing platform
  - Flexibility: on-demand and scale as you want
  - Efficiency: no need to maintain infra
- Services (relevant to this assignment)
  - Compute
    - Compute Engines: VMs
  - Big data products
    - BigQuery: Data warehouse for analytics
    - Dataproc: Hadoop and Spark
  - Storage & DBs
    - Cloud Storage: Object storage system
  - More at https://cloud.google.com/products/

GCP Setup

- Create a google account, you could use your Columbia account
- Apply for $300 credit for the first year: https://cloud.google.com/free/
- Go to Console dashboard -> Billing to check credit is there

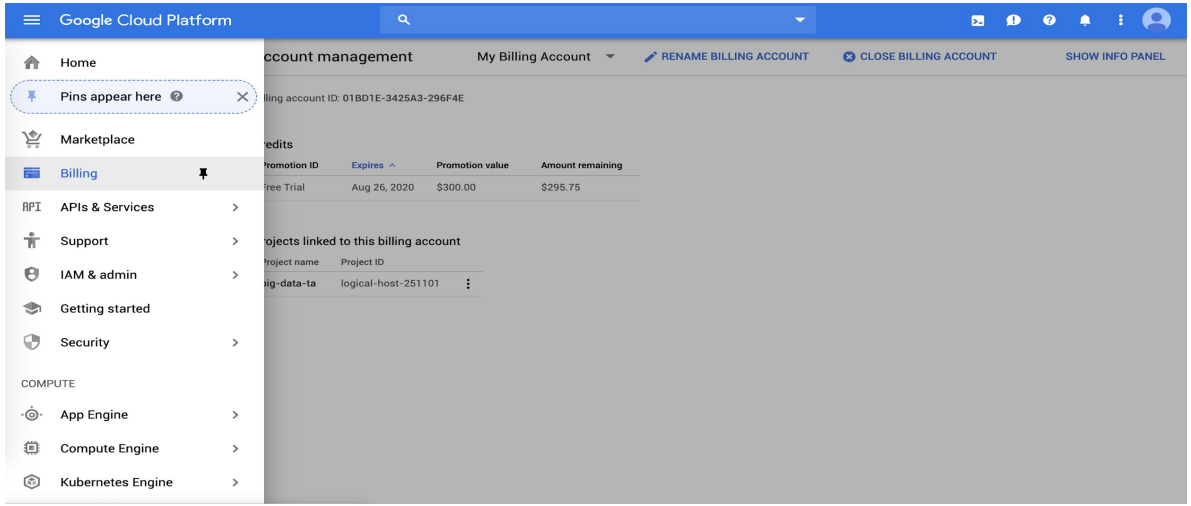

GCP: Create project

- Project: basic unit for creating, enabling, and using all GCP services
  - managing APIs, billing, permissions
  - adding and removing collaborators
- Visit console dashboard or [cloud resource manager]
- Click on “create project / new project” and complete the flow
- Ensure billing is pointing to the $300 credit

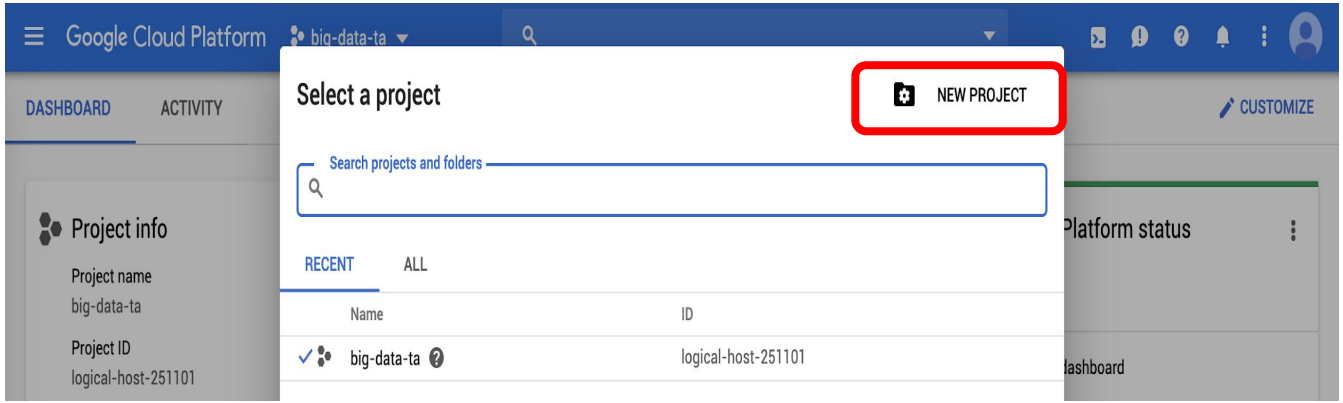

GCP: Interaction

- Graphical UI / console: Useful to create VMs, set up clusters, provision resources, manage teams, etc
- Command line tools / Cloud SDK: Useful for interacting from local host and using the resources once provisioned. E.x. ssh into instances, submit jobs, copy files, etc
- Cloud shell: Same as command line, but web-based and pre-installed with SDK and tools

GCP: console

Search for services here

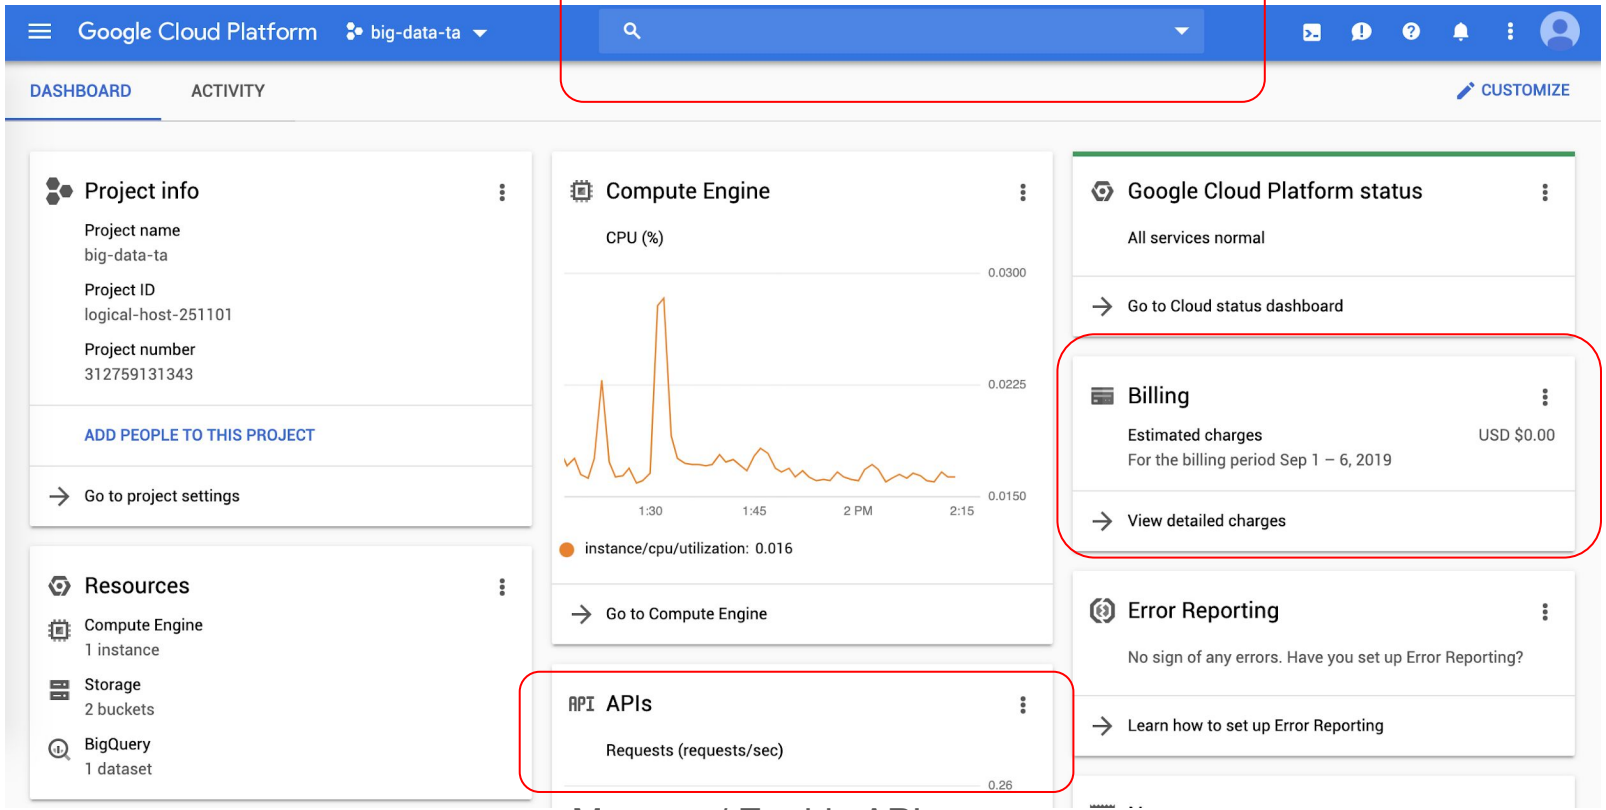

Manage / Enable APIs

GCP: Cloud SDK setup

- Install the SDK that is suitable for your local environment: https://cloud.google.com/sdk/docs/quickstarts
- Some testing after installation:
  - `gcloud auth list`
  - `gcloud components list`

```
dyn-129-236-216-148:~ frank$ gcloud components list

Your current Cloud SDK version is: 259.0.0
The latest available version is: 261.0.0
```

Components			
Status	Name	ID	Size
Update Available	BigQuery Command Line Tool	bq	< 1 MiB
Update Available	Cloud SDK Core Libraries	core	11.5 MiB
Not Installed	App Engine Go Extensions	app-engine-go	56.4 MiB
Not Installed	Cloud Bigtable Command Line Tool	cbt	7.3 MiB
Not Installed	Cloud Bigtable Emulator	bigtable	6.6 MiB
Not Installed	Cloud Datalab Command Line Tool	datalab	< 1 MiB
Not Installed	Cloud Datastore Emulator	cloud-datastore-emulator	18.4 MiB
Not Installed	Cloud Datastore Emulator (Legacy)	gcd-emulator	38.1 MiB
Not Installed	Cloud Firestore Emulator	cloud-firestore-emulator	36.8 MiB
Not Installed	Cloud Pub/Sub Emulator	pubsub-emulator	34.8 MiB
Not Installed	Cloud SQL Proxy	cloud_sql_proxy	3.7 MiB
Not Installed	Emulator Reverse Proxy	emulator-reverse-proxy	14.5 MiB
Not Installed	Google Cloud Build Local Builder	cloud-build-local	5.9 MiB
Not Installed	Google Container Registry's Docker credential helper	docker-credential-gcr	1.8 MiB
Not Installed	gcloud Alpha Commands	alpha	< 1 MiB
Not Installed	gcloud app Java Extensions	app-engine-java	85.9 MiB
Not Installed	gcloud app PHP Extensions	app-engine-php	21.9 MiB
Not Installed	gcloud app Python Extensions	app-engine-python	6.0 MiB
Not Installed	gcloud app Python Extensions (Extra Libraries)	app-engine-python-extras	28.5 MiB
Not Installed	kubectl	kubectl	< 1 MiB
Installed	Cloud Storage Command Line Tool	gsutil	3.6 MiB
Installed	gcloud Beta Commands	beta	< 1 MiB

```
To install or remove components at your current SDK version [259.0.0], run:
  $ gcloud components install COMPONENT_ID
  $ gcloud components remove COMPONENT_ID

To update your SDK installation to the latest version [261.0.0], run:
  $ gcloud components update
```

Cloud Storage

Cloud Storage

- Online file storage system
- Graphical UI through console

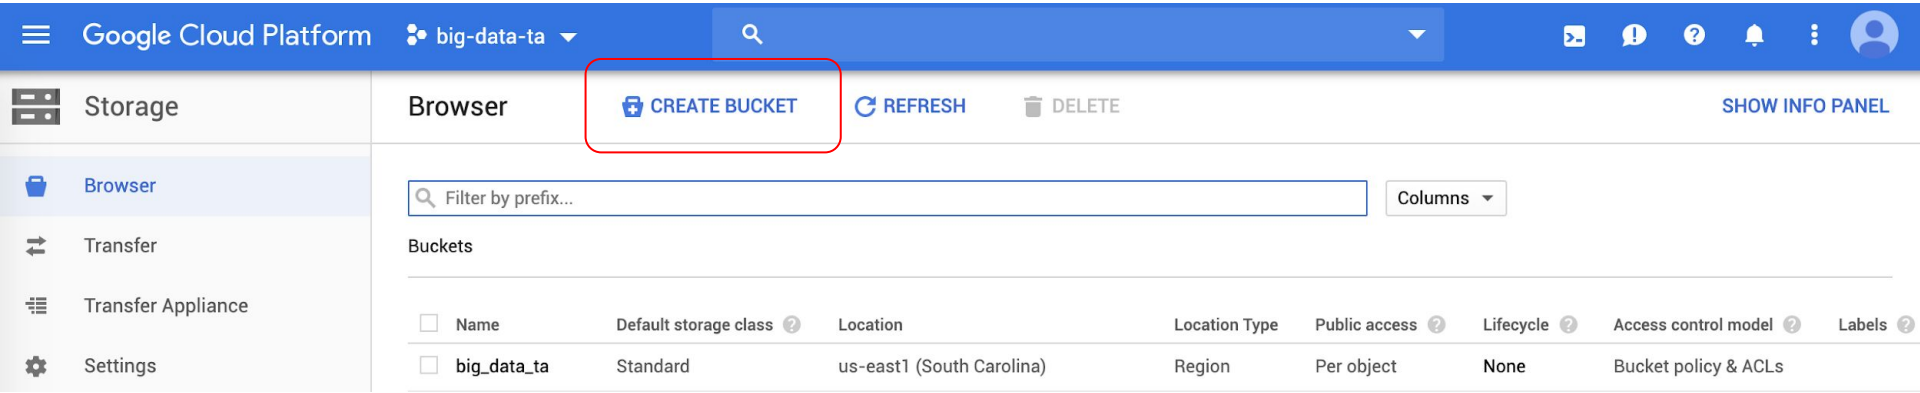

- Command line tool: `gsutil`

Cloud Storage - graphical UI

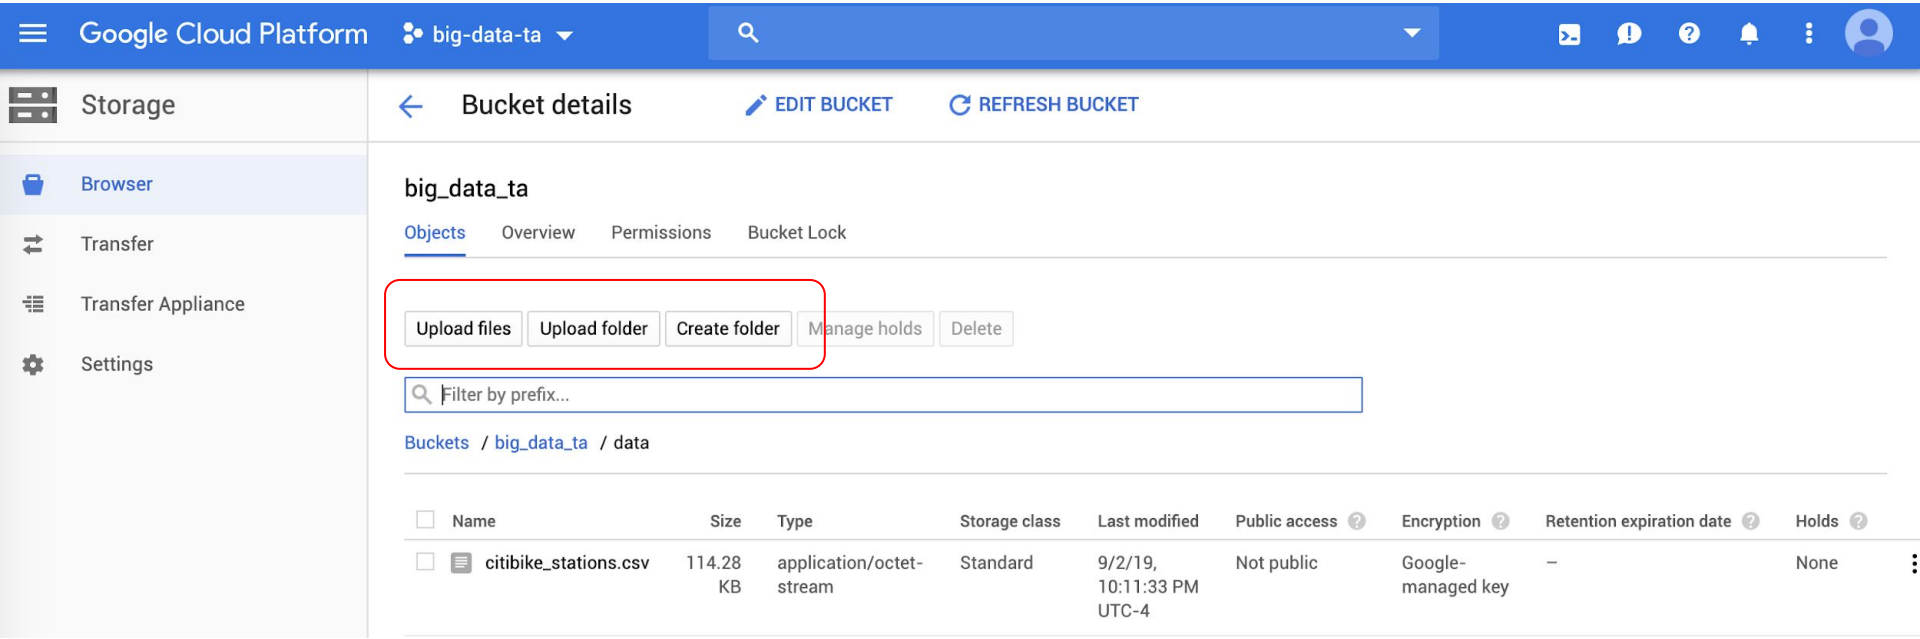

Cloud Storage - graphical UI (cont')

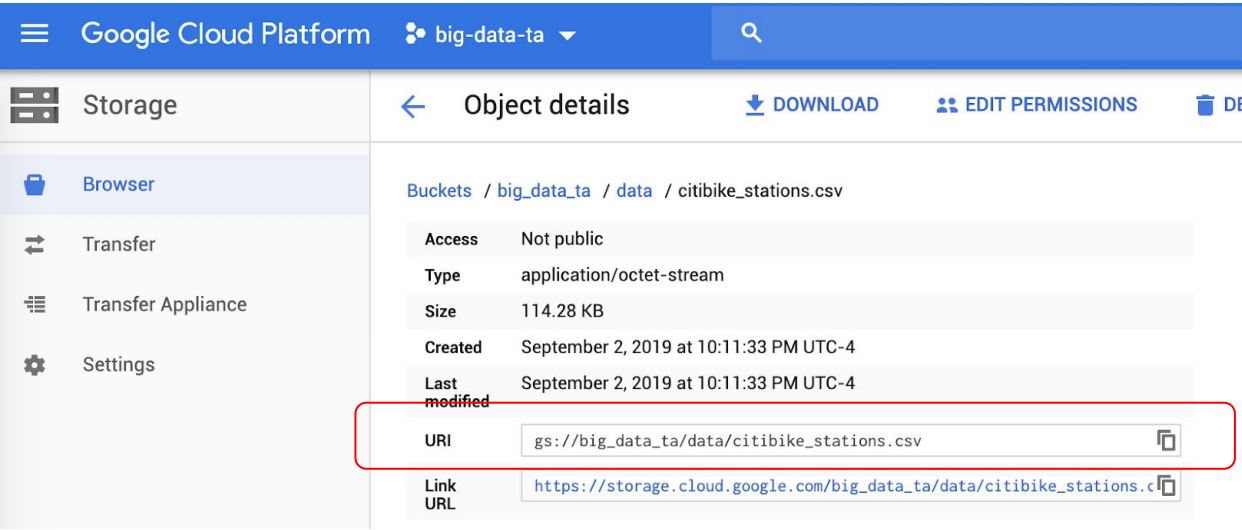

Like *filepath* on GCP,
use this in your program

Cloud Storage - `gsutil`

- Interact with Cloud Storage through command line
- Works similar to unix command line
- Useful commands:
  - Concatenate object content to stdout:
    `gsutil cat [-h] url…`
  - Copy file:
    `gsutil cp [OPTION]... src_url dst_url`
  - List files:
    `gsutil ls [OPTION]... url…`
- Explore more at https://cloud.google.com/storage/docs/gsutil

BigQuery

BigQuery

- Data warehouse for analytics
- SQL-like languages to interact with DB
- RESTful APIs for programmatic access
- Graphical UI

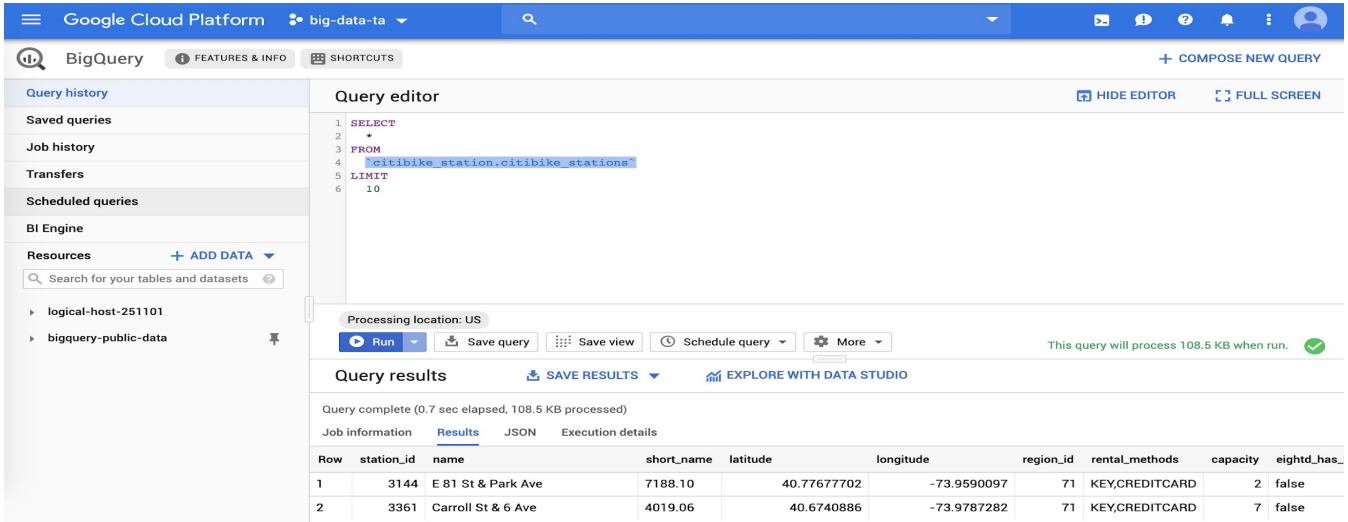

Dataproc

Dataproc

- On-demand, fully managed cloud service for running Apache Hadoop and Spark on GCP
- Cluster creation (using Cloud SDK):
  - Automatically creates VMs with Spark pre-installed
  - `gcloud dataproc clusters create <cluster-name>`
  - `gcloud beta dataproc clusters create <cluster-name> --optional-components=ANACONDA,JUPYTER --image-version=1.3 --enable-component-gateway --bucket <bucket-name> --project <project-id> --single-node --metadata 'PIP_PACKAGES=graphframes==0.6' --initialization-actions gs://dataproc-initialization-actions/python/pip-install.sh`

Install Jupyter Notebook

Cloud Storage bucket: where your jupyter notebooks are saved

Works like pip install <your package>

Dataproc - Spark execution / submit jobs

- Jupyter notebook:

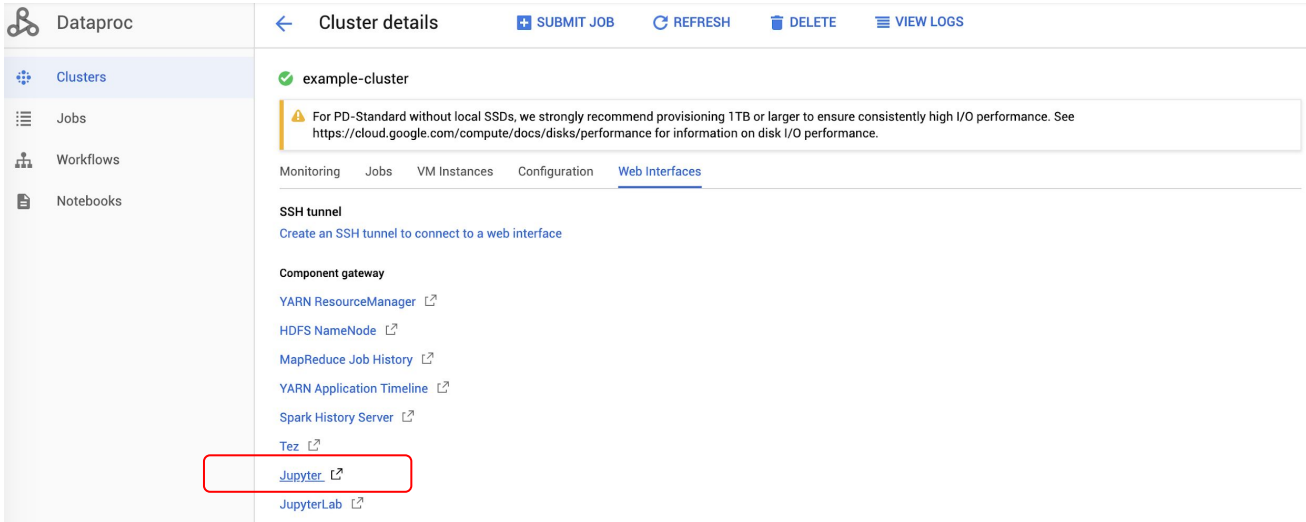

- Cloud SDK:
  - `gcloud dataproc jobs submit pyspark <your_program.py> --cluster=<cluster-name>`
  - View your jobs in console

    - program could be Cloud Storage URI or local path
    - Data should be on Cloud storage

Dataproc - Spark execution / submit jobs (cont')

- Spark shell
  - ssh into master node

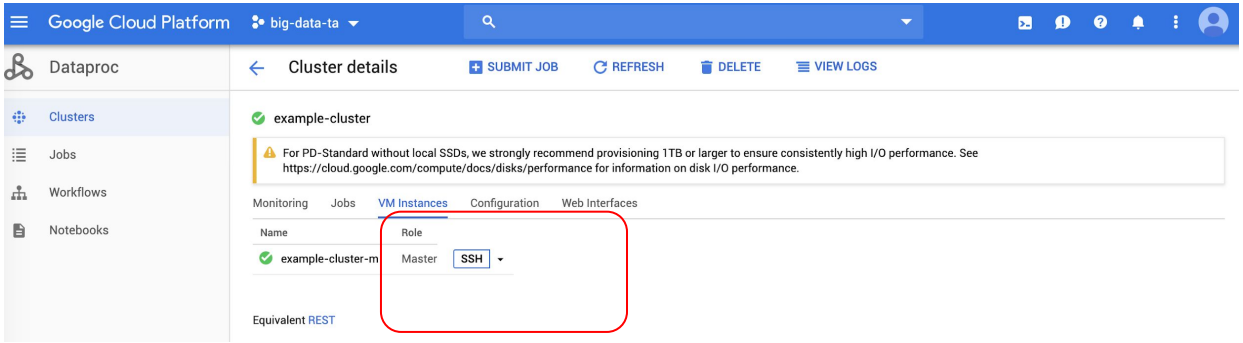

  - pyspark

```
frouyang2@example-cluster-m:~$ pyspark
Python 2.7.14 |Anaconda, Inc.| (default, Dec  7 2017, 17:05:42)
[GCC 7.2.0] on linux2
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
19/09/06 18:46:51 WARN org.apache.spark.scheduler.FairSchedulableBuilder: Fair Scheduler configuration file not fou
nd so jobs will be scheduled in FIFO order. To use fair scheduling, configure pools in fairscheduler.xml or set spa
rk.scheduler.allocation.file to a file that contains the configuration.
Welcome to
      ____              __
     / __/__  ___ _____/ /__
    _\ \/ _ \/ _ `/ __/  '_/
   /__ / .__/\_,_/_/ /_/\_\   version 2.3.3
      /_/

Using Python version 2.7.14 (default, Dec  7 2017 17:05:42)
SparkSession available as 'spark'.
>>>
```

HW0

1. Read documentations and tutorials
    a. Setup GCP and Cloud SDK
    b. Run Spark examples on Dataproc - Pi calculation and word count
    c. Familiar yourself with BigQuery
2. Two light programming questions
    a. BigQuery
    b. Spark program - Find top k most frequent words

Remember to delete your dataproc clusters when you finish executions to save money.