

EECS E6893 Big Data Analytic HW3

Chong Hu ch3467

October 31, 2019

Problem 1. Twitter data analysis with Spark Streaming

1. Hashtag result

hashtags

Schema	Details	Preview
1024	29 #iot	
1025	30 #petta	
1026	34 #bigdata	
1027	36 #viswasam	
1028	111 #bigil	
1029	114 #ai	
1030	116 #sanki.	
1031	117 #shahrukhkhan	
1032	121 #atlee	
1033	126 #sanki	
1034	235 #srk	

Figure 1: Hashtag result preview in Big Query

2. Word Count result

wordcount

Schema	Details	Preview
25	2019-10-31 10:19:00 UTC	3 good
26	2019-10-31 10:15:00 UTC	4 good
27	2019-10-31 10:17:00 UTC	5 good
28	2019-10-31 10:18:00 UTC	9 good
29	2019-10-31 10:10:00 UTC	10 movie
30	2019-10-31 10:18:00 UTC	118 movie
31	2019-10-31 10:17:00 UTC	121 movie
32	2019-10-31 10:13:00 UTC	124 movie
33	2019-10-31 10:14:00 UTC	126 movie
34	2019-10-31 10:12:00 UTC	128 movie
35	2019-10-31 10:19:00 UTC	131 movie
36	2019-10-31 10:15:00 UTC	135 movie
37	2019-10-31 10:16:00 UTC	139 movie
38	2019-10-31 10:11:00 UTC	142 movie
39	2019-10-31 10:17:00 UTC	1 spark
40	2019-10-31 10:18:00 UTC	2 spark

Figure 2: WordCount result preview in Big Query

```

.....
Time: 2019-10-31 10:31:00
.....
Waiting on bqjob_r7ecef6063fd029c6_0000016e215e688c_1 ... (0s) Current status: R
ase088c_1 ... (1s) Current status: R
Waiting on bqjob_r25ad3a09c436b116_0000016e215e7aaa_1 ... (0s) Current status: R
ase7aaa_1 ... (1s) Current status: R
Waiting on bqjob_r25ad3a09c436b116_0000016e215e7aaa_1 ...
r25ad3a09c436b116_0000016e215e7aaa_1 ... (5s) Current status: R
5) Current status: R
Job [b35241b7bb874f3680cc3c9d223d50a6] finished successfully.
DriverControlFilesUri: gs://big_data_storage/google-cloud-dataproc-metaInfo/bcc7020e-1f9b-4641-96d4-6c264c9d226b
DriverOutputResourceUri: gs://big_data_storage/google-cloud-dataproc-metaInfo/bcc7020e-1f9b-4641-96d4-6c264c9d226b
JobUuid: aa5b5d7e-ef7f-3817-8336-2e791ca6d092
Placement:
  clusterName: hw3
  clusterUuid: bcc7020e-1f9b-4641-96d4-6c264c9d226b
PySparkJob:
  mainPythonFileUri: gs://big_data_storage/google-cloud-dataproc-metaInfo/bcc7020e-1f9b-4641-96d4-6c264c9d226b
reference:
  jobId: b35241b7bb874f3680cc3c9d223d50a6
  projectId: hardy-symbol-252200
status:
  state: DONE
  stateStartTime: '2019-10-31T10:31:20.509Z'
statusHistory:
  state: PENDING
  stateStartTime: '2019-10-31T10:09:46.590Z'
  state: SETUP_DONE
  stateStartTime: '2019-10-31T10:09:46.623Z'
  details: Agent reported job success
  state: RUNNING
  stateStartTime: '2019-10-31T10:09:47.053Z'

```

Figure 3: Finish successfully in terminal

Code: sparkStreaming.py

```

1  #!/usr/bin/env python
2  # -*- coding: utf-8 -*-
3  # Columbia EECS E6893 Big Data Analytics
4  """

```

```

5  This module is the spark streaming analysis process.
6
7
8  Usage:
9      If used with dataproc:
10         gcloud dataproc jobs submit pyspark --cluster <Cluster Name> twitterHTTPClient.py
11
12         Create a dataset in BigQuery first using
13         bq mk bigdata_sparkStreaming
14
15         Remember to replace the bucket with your own bucket name
16
17
18  Todo:
19      1. hashtagCount: calculate accumulated hashtags count
20      2. wordCount: calculate word count every 60 seconds
21         the word you should track is listed below.
22      3. save the result to google BigQuery
23
24  """
25
26  import subprocess
27  import time
28
29  from pyspark import SparkConf, SparkContext
30  from pyspark.sql import SQLContext
31  from pyspark.streaming import StreamingContext
32
33  # global variables
34  bucket = "big_data_storage"
35  output_directory_hashtags = 'gs://{}/hadoop/tmp/bigquery/pyspark_output/hashtagsCount'.format(
36      bucket)
37  output_directory_wordcount = 'gs://{}/hadoop/tmp/bigquery/pyspark_output/wordcount'.format(
38      bucket)
39
40  # output table and columns name
41  output_dataset = 'twitter_analysis' # the name of your dataset in BigQuery
42  output_table_hashtags = 'hashtags'
43  columns_name_hashtags = ['hashtags', 'count']
44  output_table_wordcount = 'wordcount'
45  columns_name_wordcount = ['word', 'count', 'time']
46
47  # parameter
48  IP = 'localhost' # ip port
49  PORT = 9001 # port

```

```

50
51 STREAMTIME = 600 # time that the streaming process runs
52 # STREAMTIME = 20 # for test
53
54 WORD = ['data', 'spark', 'ai', 'movie',
55         'good'] # the words you should filter and do word count
56
57
58 # Helper functions
59 def saveToStorage(rdd, output_directory, columns_name, mode):
60     """
61     Save each RDD in this DStream to google storage
62     Args:
63     rdd: input rdd
64     output_directory: output directory in google storage
65     columns_name: columns name of dataframe
66     mode: mode = "overwrite", overwrite the file
67     mode = "append", append data to the end of file
68     """
69     if not rdd.isEmpty():
70         (rdd.toDF(columns_name)
71          .write.save(output_directory, format="json", mode=mode))
72
73
74 def saveToBigQuery(sc, output_dataset, output_table, directory):
75     """
76     Put temp streaming json files in google storage to google BigQuery
77     and clean the output files in google storage
78     """
79     files = directory + '/part-*'
80     subprocess.check_call(
81         'bq load --source_format NEWLINE_DELIMITED_JSON '
82         '--replace '
83         '--autodetect '
84         '{dataset}.{table} {files}'.format(
85             dataset=output_dataset, table=output_table, files=files
86         ).split())
87     output_path = sc._jvm.org.apache.hadoop.fs.Path(directory)
88     output_path.getFileSystem(sc._jsc.hadoopConfiguration()).delete(
89         output_path, True)
90
91
92 def hashtagCount(words):
93     """
94     Calculate the accumulated hashtags count sum from the beginning of the stream

```

```

95     and sort it by descending order of the count.
96     Ignore case sensitivity when counting the hashtags:
97         "#Ab" and "#ab" is considered to be a same hashtag
98     You have to:
99     1. Filter out the word that is hashtags.
100         Hashtag usually start with "#" and followed by a series of alphanumeric
101     2. map (hashtag) to (hashtag, 1)
102     3. sum the count of current DStream state and previous state
103     4. transform unordered DStream to a ordered Dstream
104     Hints:
105         you may use regular expression to filter the words
106         You can take a look at updateStateByKey and transform transformations
107     Args:
108         dstream(DStream): stream of real time tweets
109     Returns:
110         DStream Object with inner structure (hashtag, count)
111     """
112
113     def updateFunc(new_values, last_sum):
114         return sum(new_values) + (last_sum or 0)
115
116     hashtag = words.map(lambda x: x.lower()).filter(
117         lambda x: len(x) > 2 and x[0] == "#").map(
118         lambda x: (x, 1))
119     hashtag_cnt = hashtag.reduceByKey(lambda cnt1, cnt2: cnt1 + cnt2)
120     hashtag_cnt_total = hashtag_cnt.updateStateByKey(updateFunc)
121     return hashtag_cnt_total
122
123
124 def wordCount(words):
125     """
126     Calculate the count of 5 special words for every 60 seconds (window no overlap)
127     You can choose your own words.
128     Your should:
129     1. filter the words
130     2. count the word during a special window size
131     3. add a time related mark to the output of each window, ex: a datetime type
132     Hints:
133         You can take a look at reduceByKeyAndWindow transformation
134         Dstream is a series of rdd, each RDD in a DStream contains data from a certain interval
135         You may want to take a look of transform transformation of DStream when trying to add a time
136     Args:
137         dstream(DStream): stream of real time tweets
138     Returns:
139         DStream Object with inner structure (word, count, time)

```

```

140     """
141     word_cnt = words.map(lambda x: x.lower()).filter(lambda x: x in WORD).map(
142         lambda x: (x, 1)).reduceByKeyAndWindow(lambda x, y: x + y,
143                                                 lambda x, y: x - y, 60, 60)
144     word_cnt_total = word_cnt.transform(
145         lambda time, rdd: rdd.map(
146             lambda x: (x[0], x[1], time.strftime("%Y-%m-%d %H:%M:%S"))))
147     return word_cnt_total
148
149
150 if __name__ == '__main__':
151     # Spark settings
152     conf = SparkConf()
153     conf.setMaster('local[2]')
154     conf.setAppName("TwitterStreamApp")
155
156     # create spark context with the above configuration
157     sc = SparkContext(conf=conf)
158     sc.setLogLevel("ERROR")
159
160     # create sql context, used for saving rdd
161     sql_context = SQLContext(sc)
162
163     # create the Streaming Context from the above spark context with batch interval size 60 seconds
164     ssc = StreamingContext(sc, 60)
165     # setting a checkpoint to allow RDD recovery
166     ssc.checkpoint("~/checkpoint_TwitterApp")
167
168     # read data from port 9001
169     dataStream = ssc.socketTextStream(IP, PORT)
170     dataStream.pprint()
171
172     words = dataStream.flatMap(lambda line: line.split(" "))
173
174     # calculate the accumulated hashtags count sum from the beginning of the stream
175     topTags = hashtagCount(words)
176     topTags.pprint()
177
178     # Calculate the word count during each time period 6s
179     wordCount = wordCount(words)
180     wordCount.pprint()
181
182     # save hashtags count and word count to google storage
183     # used to save to google BigQuery
184     # You should:

```

```

185     # 1. topTags: only save the latest rdd in DStream
186     # 2. wordCount: save each rdd in DStream
187     # Hints:
188     # 1. You can take a look at foreachRDD transformation
189     # 2. You may want to use helper function saveToStorage
190     # 3. You should use save output to output_directory_hashtags, output_directory_wordcount,
191     #     and have output columns name columns_name_hashtags and columns_name_wordcount.
192
193     topTags.foreachRDD(lambda rdd: saveToStorage(rdd, output_directory_hashtags,
194                                                  columns_name_hashtags,
195                                                  mode="overwrite"))
196
197     wordCount.foreachRDD(
198         lambda rdd: saveToStorage(rdd, output_directory_wordcount,
199                                   columns_name_wordcount, mode="append"))
200
201     # start streaming process, wait for 600s and then stop.
202     ssc.start()
203     time.sleep(STREAMTIME)
204     ssc.stop(stopSparkContext=False, stopGraceFully=True)
205     # put the temp result in google storage to google BigQuery
206     saveToBigQuery(sc, output_dataset, output_table_hashtags,
207                   output_directory_hashtags)
208     saveToBigQuery(sc, output_dataset, output_table_wordcount,
209                   output_directory_wordcount)

```

twitterHTTPClient.py

```

1  #!/usr/bin/env python
2  # -*- coding: utf-8 -*-
3  # Columbia EECS E6893 Big Data Analytics
4  """
5
6  This module is used to pull data from twitter API and send data to
7  Spark Streaming process using socket. It acts like a client of
8  twitter API and a server of spark streaming. It open a listening TCP
9  server socket, and listen to any connection from TCP client. After
10 a connection established, it send streaming data to it.
11
12 Usage:
13 If used with dataproc:
14 gcloud dataproc jobs submit pyspark --cluster <Cluster Name> twitterHTTPClient.py
15
16 Make sure that you run this module before you run spark streaming process.
17 Please remember stop the job on dataproc if you no longer want to stream data.
18

```

```

19  Todo:
20      1. change the credentials to your own
21
22  """
23
24  import json
25  import socket
26
27  from tweepy import OAuthHandler
28  from tweepy import Stream
29  from tweepy.streaming import StreamListener
30
31  # credentials
32  # replace with your own credentials
33  ACCESS_TOKEN = '1186641375900647425-Yk6Mr116gbZixYb0pt2bXGiODWZiZa' # your access token
34  ACCESS_SECRET = '8xYHZbwKlF4zMdqHEWuwg01vOocHuictsQnCn4ElPT1Ny' # your access token secret
35  CONSUMER_KEY = 'Wg4xVrzE0DU8Ey4sYPQv0ntcv' # your API key
36  CONSUMER_SECRET = 'FBXUfei2Sp7W7dfuLtK8wbY5BF20KB4Ck5T5IGiE7AoEdWtQZW' # your API secret key
37
38  # the tags to track
39  tags = ['#', 'bigdata', 'spark', 'ai', 'movie']
40
41
42  class TweetsListener(StreamListener):
43      """
44      tweets listener object
45      """
46
47      def __init__(self, csocket):
48          super(TweetsListener, self).__init__()
49          self.client_socket = csocket
50
51      def on_data(self, data):
52          try:
53              msg = json.loads(data)
54              print('TEXT:{}'.format(msg['text']))
55              self.client_socket.send(msg['text'].encode('utf-8'))
56              return True
57          except BaseException as e:
58              print("Error on_data: %s" % str(e))
59              return False
60          # return True
61
62      def on_error(self, status):
63          print(status)

```



```

64         return False
65
66
67 def sendData(c_socket, tags):
68     """
69     send data to socket
70     """
71     auth = OAuthHandler(CONSUMER_KEY, CONSUMER_SECRET)
72     auth.set_access_token(ACCESS_TOKEN, ACCESS_SECRET)
73     twitter_stream = Stream(auth, TweetsListener(c_socket))
74     twitter_stream.filter(track=tags, languages=['en'])
75
76
77 class twitter_client:
78     def __init__(self, TCP_IP, TCP_PORT):
79         self.s = s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
80         self.s.bind((TCP_IP, TCP_PORT))
81
82     def run_client(self, tags):
83         try:
84             self.s.listen(1)
85             while True:
86                 print("Waiting for TCP connection...")
87                 conn, addr = self.s.accept()
88                 print("Connected... Starting getting tweets.")
89                 sendData(conn, tags)
90                 conn.close()
91         except KeyboardInterrupt:
92             exit
93
94
95 if __name__ == '__main__':
96     client = twitter_client("localhost", 9001)
97     client.run_client(tags)

```
